Cedar: Evidence Retrieval Study

1 Introduction

This document is a controlled fixture for evaluating PDF evidence retrieval. All datasets and experimental results below are fictional.

2 Method

Our method uses a sparse keyword retriever and a local reranker. Evidence passages are retained with exact page references.

3 Datasets

We evaluate on the fictional CedarQA dataset, containing 2400 training questions and 600 test questions.

4 Evaluation metrics

The evaluation metric is exact-match accuracy. The reported accuracy is 82.4 percent on CedarQA.

5 Ablation study

Removing the reranker reduces exact-match accuracy from 82.4 to 76.1 percent. This ablation suggests the reranker improves retrieval quality on this fictional benchmark.

6 Limitations

The method was evaluated only on English text. Scanned PDFs and tables were not evaluated. Generalization to other languages remains unknown.

7 Notes

No hardware latency, power consumption, clinical outcome, or deployment cost was measured. This fixture is intended for software testing only.